

2.5 — Precision and Diagnostics

ECON 480 • Econometrics • Fall 2022


Dr. Ryan Safner

Associate Professor of Economics

✉ safner@hood.edu

ryansafner/metricsF22

🌐 metricsF22.classes.ryansafner.com



Contents

Variation in $\hat{\beta}_1$

Presenting Regression Results


Diagnostics About Regression

Heteroskedasticity

Outliers

The Sampling Distribution of $\hat{\beta}_1$


$$\hat{\beta}_1 \sim N(\mathbb{E}[\hat{\beta}_1], \sigma_{\hat{\beta}_1})$$



The Sampling Distribution of $\hat{\beta}_1$

$$\hat{\beta}_1 \sim N(\mathbb{E}[\hat{\beta}_1], \sigma_{\hat{\beta}_1}^2)$$

1. **Center**¹ of the distribution: $\mathbb{E}[\hat{\beta}_1]$ (**last class**)



The Sampling Distribution of $\hat{\beta}_1$


$$\hat{\beta}_1 \sim N(\mathbb{E}[\hat{\beta}_1], \sigma_{\hat{\beta}_1}^2)$$

1. **Center**¹ of the distribution: $\mathbb{E}[\hat{\beta}_1]$ (**last class**)

2. **Precision** or **uncertainty** of the estimate (**today**)

- **Variance** $\sigma_{\hat{\beta}_1}^2$

- **Standard error**² $\sigma_{\hat{\beta}_1} = \sqrt{\text{var}(\hat{\beta}_1)}$



1. Under the 4 assumptions about u , particularly, $\text{cor}(X, u) = 0$

2. Standard “**error**” is the analog of standard *deviation* when talking about the *sampling distribution* of a sample statistic (such as \bar{X} or



The Sampling Distribution of $\hat{\beta}_1$


$$\hat{\beta}_1 \sim N(\mathbb{E}[\hat{\beta}_1], \sigma_{\hat{\beta}_1}^2)$$

1. **Center**¹ of the distribution: $\mathbb{E}[\hat{\beta}_1]$ (**last class**)

2. **Precision** or **uncertainty** of the estimate (**today**)

- **Variance** $\sigma_{\hat{\beta}_1}^2$

- **Standard error**² $\sigma_{\hat{\beta}_1} = \sqrt{\text{var}(\hat{\beta}_1)}$



1. Under the 4 assumptions about u , particularly, $\text{cor}(X, u) = 0$

2. Standard “**error**” is the analog of standard *deviation* when talking about the *sampling distribution* of a sample statistic (such as \bar{X} or



Variation in $\hat{\beta}_1$

What Affects Variation in $\hat{\beta}_1$

$$\text{var}(\hat{\beta}_1) = \frac{(SER)^2}{n \times \text{var}(X)}$$

$$\text{se}(\hat{\beta}_1) = \sqrt{\text{var}(\hat{\beta}_1)} = \frac{SER}{\sqrt{n} \times \text{sd}(X)}$$

- Variation in $\hat{\beta}_1$ is affected by 3 things:

1. Goodness of fit of the model (SER)¹]

- Larger $SER \rightarrow$ larger $\text{var}(\hat{\beta}_1)$

2. Sample size, n • Larger $n \rightarrow$ smaller $\text{var}(\hat{\beta}_1)$

3. Variance of X

- Larger $\text{var}(X) \rightarrow$ smaller $\text{var}(\hat{\beta}_1)$


1. ^{magenta[†]} Recall from last class, the Standard Error of the Regression $\hat{\sigma}_u = \sqrt{\frac{\sum \hat{u}_i^2}{n-2}}$



Variation in $\hat{\beta}_1$: Goodness of Fit


Model With Better Fit

Lower SER lowers variation in $\hat{\beta}_1$



Model With Worse Fit


Higher SER raises variation in $\hat{\beta}_1$



Variation in $\hat{\beta}_1$: Sample Size


Model With Fewer Observations

Smaller n raises variation in $\hat{\beta}_1$



Model With More Observations


Larger n lowers variation in $\hat{\beta}_1$



Variation in $\hat{\beta}_1$: Variation in X


Model With More Variation in X

Larger $\text{var}(X)$ lowers variation in $\hat{\beta}_1$



Model With Less Variation in X

Smaller $\text{var}(X)$ raises variation in $\hat{\beta}_1$



Presenting Regression Results

Our Class Size Regression

```
1 school_reg %>% summary()
```

Call:

```
lm(formula = testscr ~ str, data = ca_school)
```

Residuals:

Min	1Q	Median	3Q	Max
-47.727	-14.251	0.483	12.822	48.540

Coefficients:

	Estimate	Std. Error	t value	Pr(> t)
(Intercept)	698.9330	9.4675	73.825	< 2e-16 ***
str	-2.2798	0.4798	-4.751	2.78e-06 ***

Signif. codes: 0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1

- How can we present all of this information in a tidy way?



Our Class Size Regression

```
1 library(broom)
2 school_reg %>% tidy()
```

term	estimate	std.error	statistic	p.value
(Intercept)	698.932952	9.4674914	73.824514	0.0e+00
str	-2.279808	0.4798256	-4.751327	2.8e-06

```
1 school_reg %>% glance()
```

r.squared	adj.r.squared	sigma	statistic	p.value	df	logLik	AIC	BIC	deviance
0.0512401	0.0489703	18.58097	22.57511	2.8e-06	1	-1822.25	3650.499	3662.62	1822.25

- Better (?), but still not how you see regressions reported in reports...especially when you have many regression models!



Regression Tables

- Professional journals and papers often have a **regression table**, including:
 - Estimates of $\hat{\beta}_0$ and $\hat{\beta}_1$
 - Standard errors of $\hat{\beta}_0$ and $\hat{\beta}_1$ (often below, in parentheses)
 - Indications of statistical significance (often with asterisks)
 - Measures of regression fit: R^2 , SER , etc
- Later: multiple rows & columns for multiple variables & models

	Test Score
Constant	698.93***
	(9.47)
STR	-2.28***
	(0.48)
n	420
R ²	0.05
SER	18.54
* p < 0.1, ** p < 0.05, *** p < 0.01	



Regression Output Tables

- A number of packages (and documentation/guides) that will make nice regression output tables for you:
 - `modelsummary`
 - `stargazer` (and a good `cheat sheet`)
 - `huxtable`

	Test Score
Constant	698.93***
	(9.47)
STR	-2.28***
	(0.48)
n	420
R ²	0.05
SER	18.54
* p < 0.1, ** p < 0.05, *** p < 0.01	



Using `modelsummary` I

- You will need to first
`install.packages("modelsummary")`
- Load with `library(modelsummary)`
- Command: `modelsummary()`
- Main argument is the name of your `lm` regression object
- Default output is *fine*, but often we want to customize a bit!

```
1 # install.packages("modelsummary") # install first
2 # load package
3 library(modelsummary)
4
5 modelsummary(school_reg) # our regression
```

	Model 1
(Intercept)	698.933
	(9.467)
str	-2.280
	(0.480)
Num.Obs.	420
R2	0.051
R2 Adj.	0.049
AIC	3650.5
BIC	3662.6
F	22.575
RMSE	18.54



Using `modelsummary` II

- Whole command is `modelsummary()`, everything will go in `()`
1. `models`, a `list()` of models to use, can give a name to each model, will show up as column title in table

```
1 models = list("Test Score" = school_reg) # set name to "Test Score"
```

2. `coef_rename` if you want to rename any independent variables as something nicer than their names in the dataset

- `"old name" = "new name"` (yes annoying!)

```
1 coef_rename = list("(Intercept)" = "Constant",  
2                     "str" = "Student Teacher Ratio")
```



Using `modelsummary` III

- Whole command is `modelsummary()`, everything will go in `()`
3. `gof_map`: a `list()` of goodness of fit statistics, can customize what you want to include/exclude, what you want to label them in the table...a bit advanced, here's what I like:

```
1 gof_map = list(
2   list("raw" = "nobs", "clean" = "n", "fmt" = 0),
3   list("raw" = "r.squared", "clean" = "R<sup>2</sup>", "fmt" = 2),
4   #list("raw" = "adj.r.squared", "clean" = "Adj. R<sup>2</sup>", "fmt" = 2), # we'll want this later!
5   list("raw" = "rmse", "clean" = "SER", "fmt" = 2)
6 )
```

4. Other minor options (combine with commas):

```
1 fmt = 2, # round to 2 decimals
2 output = "html" # depending on type of document creating; pdf would be "latex"
3 escape = FALSE # allows formatting of things like <sup>2</sup>
4 stars = c('*' = .1, '**' = .05, '***' = 0.01) # show significance levels if set to true, I don't like the d
```



Using modelsummary IV

```

1 modelsummary(models = list("Test Score" = school_reg),
2               fmt = 2, # round to 2 decimals
3               output = "html",
4               coef_rename = c("(Intercept)" = "Constant",
5                               "str" = "STR"),
6               gof_map = list(
7                 list("raw" = "nobs", "clean" = "n", "fmt" = 0)
8                 list("raw" = "r.squared", "clean" = "R<sup>2</sup>"),
9                 #list("raw" = "adj.r.squared", "clean" = "Adj.
10                 list("raw" = "rmse", "clean" = "SER", "fmt" =
11                 ),
12               escape = FALSE,
13               stars = c('*' = .1, '**' = .05, '***' = 0.01)
14 )

```


	Test Score
Constant	698.93***
	(9.47)
STR	-2.28***
	(0.48)
n	420
R ²	0.05
SER	18.54
* p < 0.1, ** p < 0.05, *** p < 0.01	



modelplot() in modelsummary

Also nice about the `modelsummary` package is the command `modelplot()`


```
1 modelplot(school_reg)
```



modelplot() in modelsummary

Also nice about the `modelsummary` package is the command `modelplot()`

```
1 modelplot(school_reg,  
2           coef_omit = 'Intercept') # don't show intercept
```




Though You Could Make It Yourself in **ggplot**

- Use the `conf.low` and `conf.high` (from a `tidy` regression) as `xmin` and `xmax` aesthetics inside `geom_errorbarh()`.

```

1 school_reg %>%
2   tidy(conf.int = TRUE) %>%
3   filter(term == "str") %>%
4   ggplot()+
5   aes(x = estimate,
6       y = term)+
7   geom_point(size = 3)+
8   geom_errorbarh(aes(xmin = conf.low, xmax = conf.
9                     height = 0.1)+ # height of whisker
10  theme_light()
```



Diagnostics About Regression

Diagnostics: Residuals I

- We often look at the residuals of a regression to get more insight about its **goodness of fit** and its **bias**
- Recall `broom`'s `augment` creates some useful new variables
 - `.fitted` are fitted (predicted) values from model, i.e. \hat{Y}_i
 - `.resid` are residuals (errors) from model, i.e. \hat{u}_i



Diagnostics: Residuals II

- Often a good idea to store in a new object (so we can make some plots)

```
1 aug_reg <- augment(school_reg)
2
3 aug_reg %>% head()
```

testscr	str	.fitted	.resid	.hat	.sigma	.cooks	.std.resid
690.80	17.88991	658.1474	32.65260	0.0044244	18.53408	0.0068925	1.7612148
661.20	21.52466	649.8608	11.33917	0.0047485	18.59490	0.0008927	0.6117112
643.60	18.69723	656.3069	-12.70689	0.0029742	18.59279	0.0006996	-0.6848850
647.70	17.35714	659.3620	-11.66198	0.0058575	18.59441	0.0011673	-0.6294767
640.85	18.67133	656.3659	-15.51592	0.0030072	18.58766	0.0010548	-0.8363024
605.55	21.40625	650.1308	-44.58076	0.0044603	18.47411	0.0129531	-2.4046387



Recall: Assumptions about Errors

- We make **4 critical assumptions about u** :

1. The expected value of the errors is 0

$$\mathbb{E}[u] = 0$$

2. The variance of the errors over X is constant:

$$\text{var}(u|X) = \sigma_u^2$$

3. Errors are not correlated across observations:

$$\text{cor}(u_i, u_j) = 0 \quad \forall i \neq j$$

4. There is no correlation between X and the error term:

$$\text{cor}(X, u) = 0 \text{ or } E[u|X] = 0$$



Assumptions 1 and 2: Errors are i.i.d.

- Assumptions 1 and 2 assume that errors are coming from the same (*normal*) distribution

$$u \sim N(0, \sigma_u)$$


- Assumption 1: $E[u] = 0$
- Assumption 2: $sd(u|X) = \sigma_u$
 - virtually always unknown...
- We often can visually check by plotting a **histogram** of u



Plotting a Histogram of Residuals

Plot

Code



Checking the Distribution of Residuals

```
1 school_reg %>% summary()
```

Call:

```
lm(formula = testscr ~ str, data = ca_school)
```

Residuals:

Min	1Q	Median	3Q	Max
-47.727	-14.251	0.483	12.822	48.540

Coefficients:

	Estimate	Std. Error	t value	Pr(> t)
(Intercept)	698.9330	9.4675	73.825	< 2e-16 ***
str	-2.2798	0.4798	-4.751	2.78e-06 ***

Signif. codes: 0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1

```
1 aug_reg %>%
2   summarize(E_u = mean(.resid),
3             sd_u = sd(.resid))
```

E_u	sd_u
-----	------

0	18.55878
---	----------



Residual Plot

- We often plot a **residual plot** to see any odd patterns about residuals
 - x -axis are \hat{Y}_i values (`.fitted`)
 - y -axis are u_i values (`.resid`)

Plot	Code
------	------



Heteroskedasticity

Homoskedasticity

- “**Homoskedasticity**,” variance of the residuals over X is constant, written:

$$\text{var}(u|X) = \sigma_u^2$$

- Knowing the value of X does not affect the variance (spread) of the errors



Heteroskedasticity I

- “**Heteroskedasticity**,” variance of the residuals over X is **NOT** constant:

$$\text{var}(u|X) \neq \sigma_u^2$$

- **This does not cause $\hat{\beta}_1$ to be biased**, but it does cause the standard error of $\hat{\beta}_1$ to be incorrect
- This **does** cause a problem for **inference**!
 - Specifically, it will make $se(\hat{\beta}_1)$ wrong (often too small)¹



Heteroskedasticity II

- Recall the formula for the standard error of $\hat{\beta}_1$:

$$se(\hat{\beta}_1) = \sqrt{var(\hat{\beta}_1)} = \frac{SER}{\sqrt{n} \times sd(X)}$$

- This *assumes* homoskedasticity (Assumption 2)



Heteroskedasticity III

- A better formula for estimating standard errors that are **robust** to heteroskedasticity (called .hi[“robust standard errors”]):


$$se(\hat{\beta}_1) = \sqrt{\frac{\sum_{i=1}^n (X_i - \bar{X})^2 \hat{u}^2}{\left[\sum_{i=1}^n (X_i - \bar{X})^2 \right]^2}}$$

- Don't learn formula, **do learn what heteroskedasticity is and how it affects our model!**




Visualizing Heteroskedasticity I

- Our original scatterplot with regression line
- Does the spread of the errors change over different values of str ?
 - No: homoskedastic
 - Yes: heteroskedastic



Visualizing Heteroskedasticity




- Notice the distribution of \hat{u} , changes for different values of STR, and $\sigma_{\hat{u}}$ is not constant




More Obvious Heteroskedasticity

- Visual cue: data is “fan-shaped”
 - Data points are closer to line in some areas
 - Data points are more spread from line in other areas




More Obvious Heteroskedasticity




What Might Cause Heteroskedastic Errors?

$$wage_i = \beta_0 + \beta_1 educ_i + u_i$$



What Might Cause Heteroskedastic Errors?

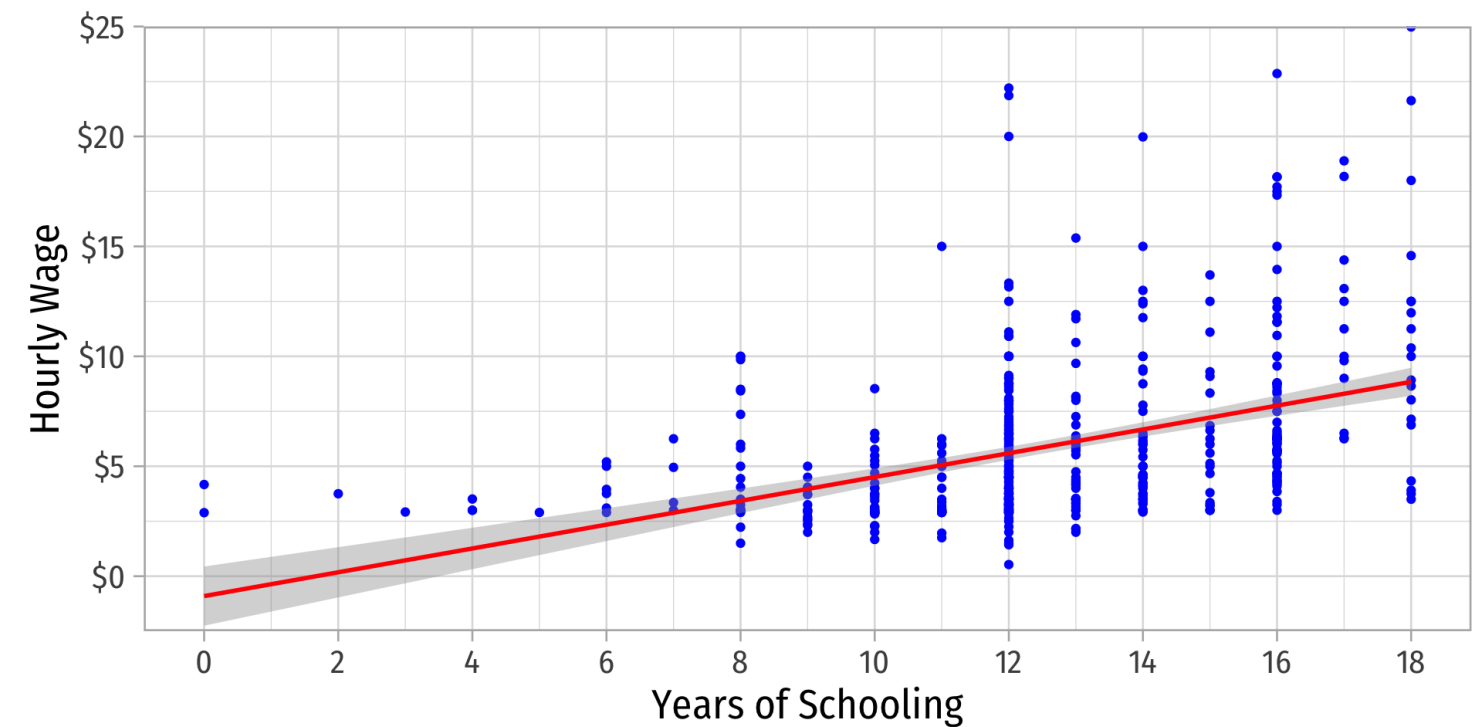
$$wage_i = \beta_0 + \beta_1 educ_i + u_i$$



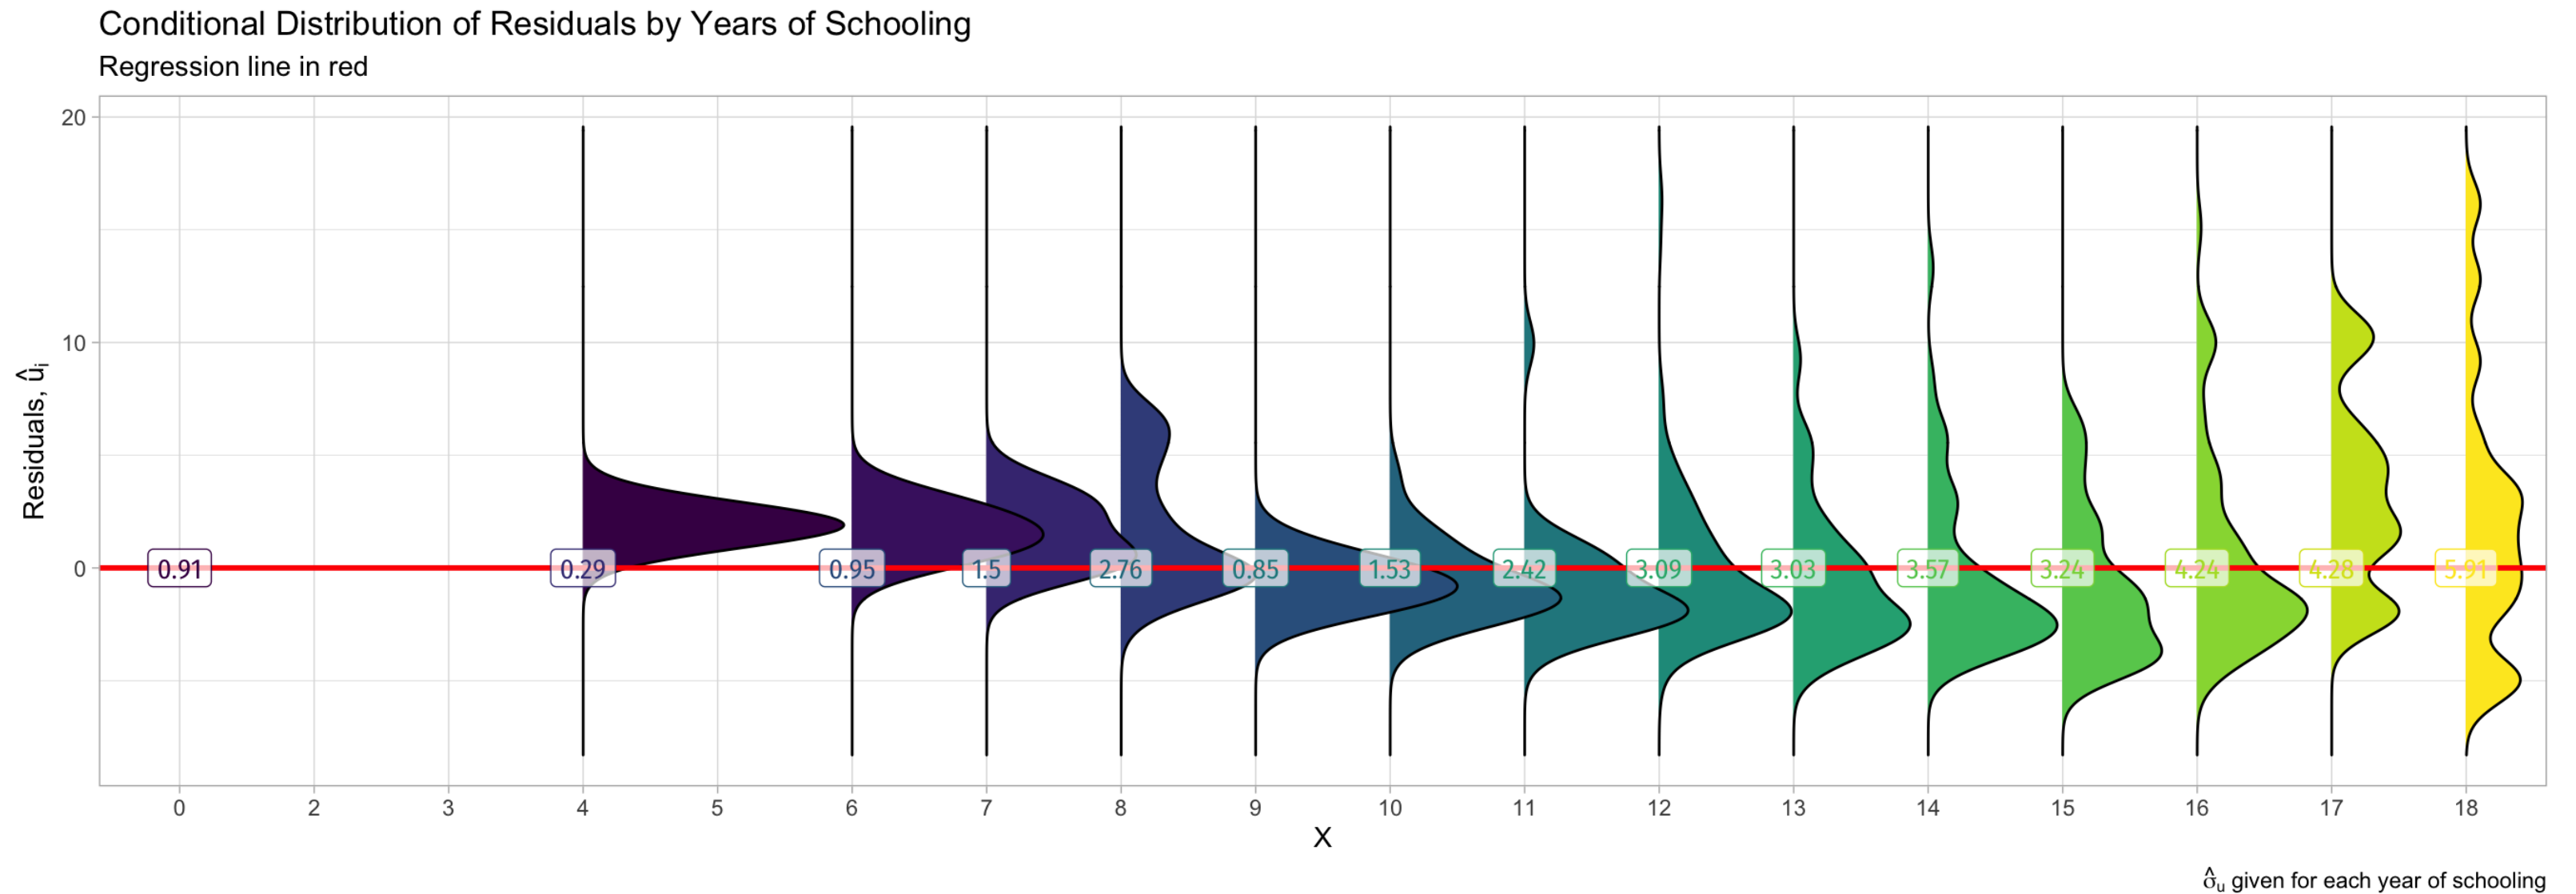
What Might Cause Heteroskedastic Errors?

$$wage_i = \beta_0 + \beta_1 educ_i + u_i$$

	Wage
Intercept	-0.90 (0.68)
Years of Schooling	0.54*** (0.05)
n	526
R ²	0.16
SER	3.37
* p < 0.1, ** p < 0.05, *** p < 0.01	



What Might Cause Heteroskedastic Errors?



Detecting Heteroskedasticity I

- Several tests to check if data is heteroskedastic
- One common test is **Breusch-Pagan test**
- Can use the `lmtest` package's function `bptest()`
 - H_0 : homoskedastic¹
 - If $p\text{-value} < 0.05$, reject $H_0 \implies$ heteroskedastic

```
1 library("lmtest")
2 school_reg %>% bptest()
```

studentized Breusch-Pagan test

data: .

BP = 5.7936, df = 1, p-value = 0.01608

- Since $p < 0.05$, can reject H_0 that errors are homoskedastic and conclude they are *heteroskedastic*



How About the Wages Regression?

```
1 wage_reg %>% bptest()
```

studentized Breusch-Pagan test

data: .

BP = 15.306, df = 1, p-value = 9.144e-05



Fixing Heteroskedasticity I

- Heteroskedasticity is easy to fix with software that can calculate **robust standard errors** (using the more complicated formula above)
- Easiest method is to use `estimat`r package
 - `lm_robust()` command (*instead of* `lm`) to run regression
 - set `se_type = "stata"` to calculate robust SEs using the formula above¹

```
1 #install.packages("estimat")
2 library(estimat)
```



Fixing Heteroskedasticity II

```
1 school_reg_robust <- lm_robust(testscr ~ str, data = ca_school,
2                               se_type = "stata")
3
4 school_reg_robust
```

	Estimate	Std. Error	t value	Pr(> t)	CI Lower	CI Upper
(Intercept)	698.932952	10.3643599	67.436191	9.486678e-227	678.560192	719.305713
str	-2.279808	0.5194892	-4.388557	1.446737e-05	-3.300945	-1.258671

	DF
(Intercept)	418
str	418

```
1 school_reg_robust %>% summary()
```

Call:

```
lm_robust(formula = testscr ~ str, data = ca_school, se_type = "stata")
```

Standard error type: HC1

Coefficients:

	Estimate	Std. Error	t value	Pr(> t)	CI Lower	CI Upper	DF
(Intercept)	698.93	10.3644	67.436	9.487e-227	678.560	719.306	418
str	-2.28	0.5195	-4.389	1.447e-05	-3.301	-1.259	418

Multiple R-squared: 0.05124 , Adjusted R-squared: 0.04897

F-statistic: 19.26 on 1 and 418 DF, p-value: 1.447e-05



Fixing Heteroskedasticity III

```
1 # can tidy, glance, augment, etc
2 school_reg_robust %>% tidy()
```

term	estimate	std.error	statistic	p.value	conf.low	conf.high	df	ou
(Intercept)	698.932952	10.3643599	67.436191	0.00e+00	678.560192	719.305713	418	tes
str	-2.279808	0.5194892	-4.388557	1.45e-05	-3.300945	-1.258671	418	tes

```
1 school_reg_robust %>% glance()
```

r.squared	adj.r.squared	statistic	p.value	df.residual	nobs	se_type
0.0512401	0.0489703	19.25943	1.45e-05	418	420	HC1



Showing The Effect of Heteroskedasticity (on $se(\hat{\beta}_1)$)

```
1 modelsummary(models = list("Normal SE" = school_re
2                               "Robust SE" = school_re
3                               fmt = 2, # round to 2 decimals
4                               output = "html",
5                               coef_rename = c("(Intercept)" = "Cons
6                                               "str" = "STR"),
7                               gof_map = list(
8                                   list("raw" = "nobs", "clean" = "n",
9                                   list("raw" = "r.squared", "clean" =
10                                   #list("raw" = "adj.r.squared", "cle
11                                   list("raw" = "rmse", "clean" = "SER
12                               ),
13                               escape = FALSE,
14                               stars = c('*' = .1, '**' = .05, '***'
15 )
```

	Normal SE	Robust SE
Constant	698.93***	698.93***
	(9.47)	(10.36)
STR	-2.28***	-2.28***
	(0.48)	(0.52)
n	420	420
R ²	0.05	0.05
SER	18.54	18.54
* p < 0.1, ** p < 0.05, *** p < 0.01		

- What changed?




Outliers

Outliers Can Bias OLS! I


- **Outliers** can affect the slope (and intercept) of the line and add **bias**
 - May be result of human error (measurement, transcribing, etc)
 - May be meaningful and accurate
- In any case, compare how including/dropping outliers affects regression and always discuss outliers!



Outliers Can Bias OLS! II




Outliers Can Bias OLS! II



Outliers Can Bias OLS! III

	Original	With Outliers
Constant	698.93***	641.40***
	(9.47)	(11.21)
STR	-2.28***	0.71
	(0.48)	(0.57)
n	420	423
R ²	0.05	0.00
SER	18.54	23.71
* p < 0.1, ** p < 0.05, *** p < 0.01		



Detecting Outliers

- The `car` package has an `outlierTest` command to run on the regression

```
1 # install.packages("car")
2 library("car")
3
4 # Use Bonferonni test
5 outlierTest(school_outlier_reg) # will point out which obs #s seem outliers
```

```
      rstudent unadjusted p-value Bonferroni p
422 8.822768      3.0261e-17  1.2800e-14
423 7.233470      2.2493e-12  9.5147e-10
421 6.232045      1.1209e-09  4.7414e-07
```

```
1 # find these observations
2 ca_school_outliers %>%
3   slice(c(422,423,421)) # find observations 422, 423, 421
```

observat	district	testscr	str
422	Crazy District 2	850	28
423	Crazy District 3	820	29
421	Crazy District 1	800	30

